


Interprocess Communication



From **Coulouris, Dollimore, Kindberg and Blair**
Distributed Systems:
Concepts and Design


Edition 5, © Addison-Wesley 2012

The characteristics of interprocess communication

- Messages are sent to (Internet address, local port) pairs.
- A port has exactly one receiver but can have many senders
- Processes may use multiple ports to receive messages
- Validity: A point-to-point message service is reliable if messages are guaranteed to be delivered despite packet drop
- Integrity: Messages must arrive uncorrupted and without duplication
- Ordering: messages be delivered in sender order

Application and middleware layers use the services provided by the network and transport layers through socket API.

Sockets



Socket programming

Goal: learn how to build client/server application that communicate using sockets

Socket API

- introduced in BSD4.1 UNIX, 1981
- explicitly created, used, released by apps
- client/server paradigm
- two types of transport service via socket API:
- unreliable datagram
- reliable, byte stream-oriented

socket

a *host-local, application-created, OS-controlled* interface (a "door") into which application process can *both send and receive* messages to/from another application process

Processes-to-process communication

Process: program running within a host.

within same host, two processes communicate using **inter-process communication** (shared memory defined by OS).

processes in different hosts communicate by exchanging **messages** using transport layer

Client process: process that initiates communication

Server process: process that waits to be contacted

- ❑ Note: applications with P2P architectures have client processes & server processes

Addressing processes

to receive messages,
process must have
identifier

host device has unique 32-bit IP address

Q: does IP address of host on which process runs suffice for identifying the process?

A: No, *many* processes can be running on the same host

identifier includes both **IP address** and **port number** associated with the process

What is a port number?

16 bits integer used by transport layer to identify end points (processes) on a host

well-known ports: 1 – 1023 Telnet 23; FTP 21; HTTP 80

registered ports: 1024 – 49151

dynamic or private ports: 49152 - 65535



To communicate, client must know the server's **IP address**, and **port number**.
How will the server know the client's IP address and port number?

Sockets

API, an interface, gate, door
between a process and
transport layer

A socket must be bound to a
local port

Is (IP addr, port) enough to
identify a socket?



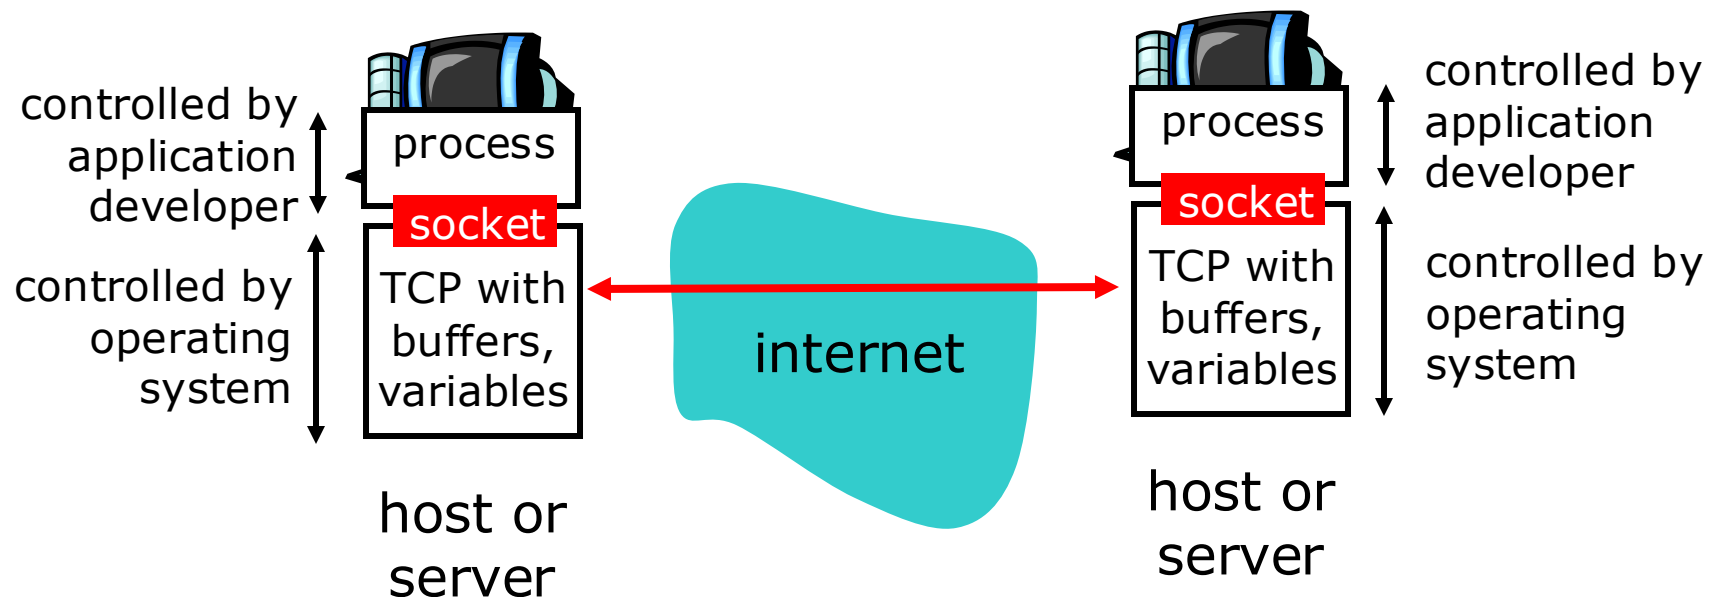
Internet address = 138.37.94.248

Internet address = 138.37.88.249

Socket-programming using TCP

Socket: a door between application process and end-end-transport protocol (UCP or TCP)

TCP service: reliable transfer of **bytes** from one process to another



Socket programming *with TCP*

Client must contact server

- server process must first be running
- server must have created socket (door) that welcomes client's contact

Client contacts server by:

- creating client-local TCP socket
- specifying IP address, port number of server process
- When **client creates socket**: client TCP establishes connection to server TCP

When contacted by client, **server**

TCP creates new socket for server process to communicate with client

- allows server to talk with multiple clients
- source port numbers used to distinguish clients

application viewpoint

 *TCP provides reliable, in-order transfer of bytes ("pipe") between client and server*

Socket with TCP

Web server has two sockets opened: one for each web page it is serving. These sockets are differentiated by the destination port numbers.

Server Sockets		Socket 1	Socket 2	Socket 3
Source port numbers assigned by local host →	Transport	TCP	TCP	TCP
	Source Port	80	80	25
	Source IP Addr	192.168.1.102	192.168.1.102	192.168.1.102
Destination port numbers assigned by remote host →	Destination Port	10378	24543	43876
	Destination IP Addr	192.168.1.101	192.168.1.101	192.168.1.101
		↑ Web server web page X	↑ Web server web page Y	↑ Email server

Client/server socket interaction: TCP

Server (stand-by, waiting for requests)

Client (initiate the request)

1 Create socket
(Claim resources/ available phone)

4 Create socket
(Claim resources/ available phone)

2 Bind port
(Claim ID on this machine/
get a phone extension No.)

-
(Don't care about public port/ phone #)

3 Listen/ **accept**
(Wait for connections/ Wait for phone call)

5 **Connect**
(connect to server/ call others)

6 **Send/ receive**
(Communication/ Chat on phone)

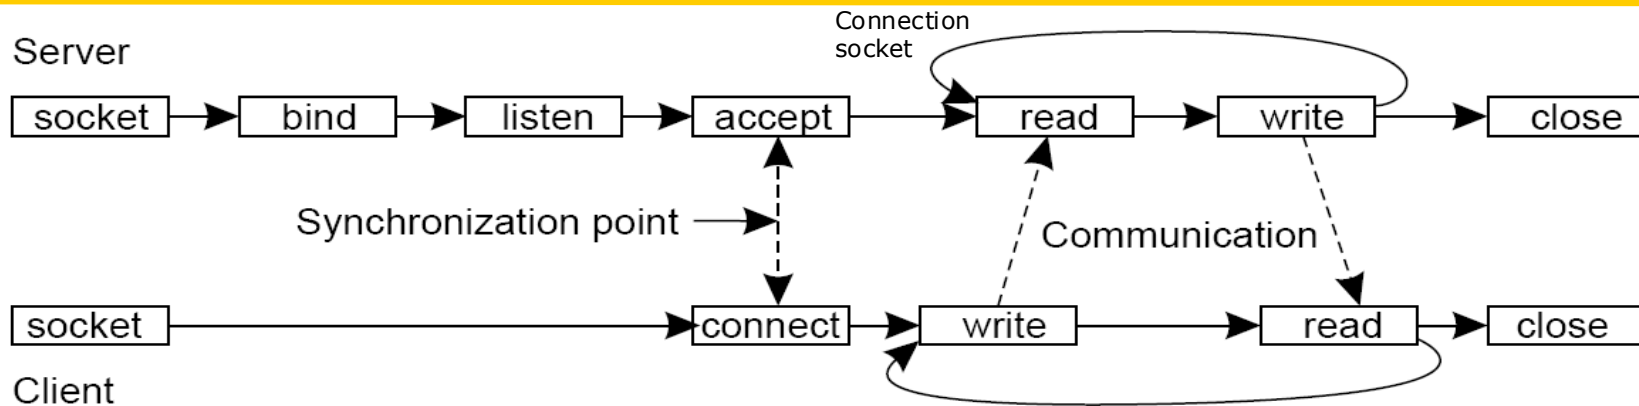
6 **Send/ receive**
(Communication/ Chat on phone)

7/8 Close socket
(End communication/ Hang up the phone)

7/8 Close socket
(End communication/ Hang up the phone)

Red word: wait for the other side

TCP Socket Primitives



<u>Primitive</u>	<u>Function</u>
Socket	Create a new communication endpoint
Bind	Attach a local address to a socket
Listen	Announce willingness to accept connections
Accept	Block caller until a connection request arrives
Connect	Actively attempt to establish a connection
Send	Send some data over the connection
Recv	Receive some data over the connection
Close	Release the connection


Client/server socket interaction: TCP

Server (stand-by, waiting for requests) **Client** (initiate the request)

Create socket
(Claim resources/ available phone)

Create socket
(Claim resources/ available phone)

Bind port
(Claim ID on this machine/
get a phone extension No.)



Connect
(connect to server/ call others)



Send/ receive
(Communication/ Chat on phone)

Send/ receive
(Communication/ Chat on phone)

Close socket
(End communication/ Hang up the phone)

Close socket
(End communication/ Hang up the phone)

Client/server socket interaction: TCP

Server (running on `hostid`, port `x`)

create socket,
port=`x`, for
incoming request:
`welcomeSocket =`
`ServerSocket()`

wait for incoming
connection request
`connectionSocket =`
`welcomeSocket.accept()`

read request from
`connectionSocket`

write reply to
`connectionSocket`

close
`connectionSocket`

Client (running on hostname `?`, port `?`)

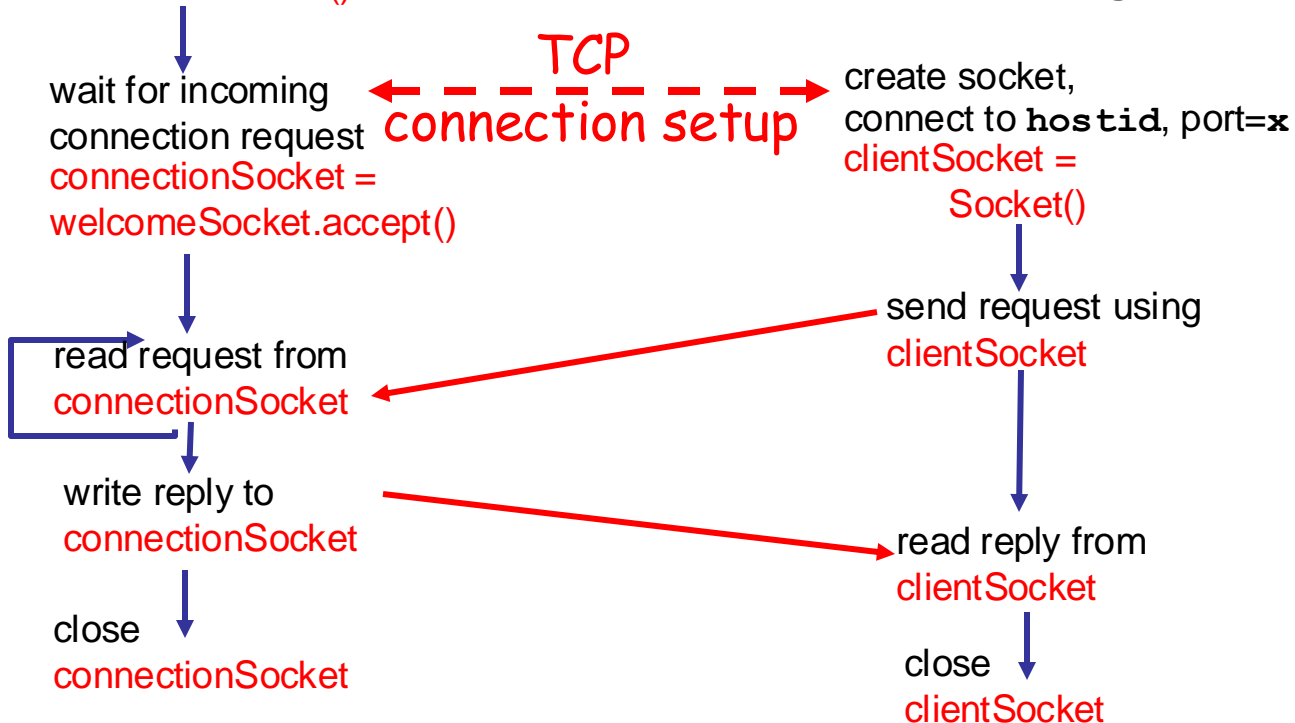
create socket,
connect to `hostid`, port=`x`
`clientSocket =`
`Socket()`

send request using
`clientSocket`

read reply from
`clientSocket`

close
`clientSocket`

TCP
connection setup



Connection-oriented TCP

❑ TCP socket identified by 4-tuple:

- source IP address
- source port number
- destion IP address
- destion port number

❑ Receive host uses all four values to direct segment to appropriate socket

Socket programming *with UDP*

UDP: no “connection” between client and server

- no handshaking
- sender explicitly attaches IP address and port of destination
- server must extract IP address, port of sender from received datagram

UDP: transmitted data may be received out of order, or lost

application viewpoint

UDP provides unreliable transfer of groups of bytes (“datagrams”) between client and server

Socket programming *with UDP*

Server (running on `hostid`)

create socket,
port=`x`, for
incoming request:
`serverSocket =`
`DatagramSocket()`

read request from
`serverSocket`

write reply to
`serverSocket`
specifying client
host address,
port number

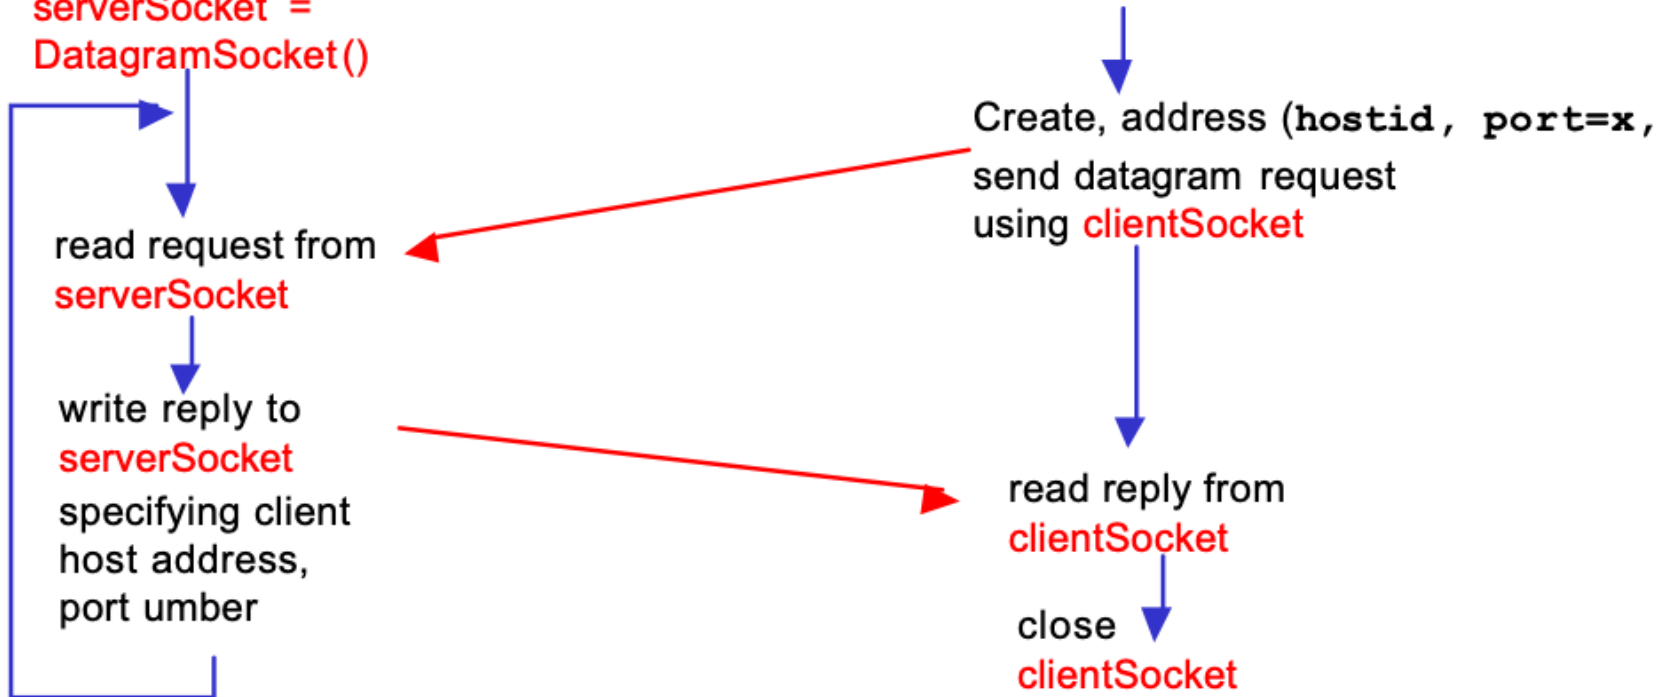
Client

create socket,
`clientSocket =`
`DatagramSocket()`

Create, address (`hostid`, `port=x`,
send datagram request
using `clientSocket`


read reply from
`clientSocket`

close
`clientSocket`



Multicast Communication At Network Layer

Unicast Communication



In unicast routing, the router forwards the received packet through only one of its interfaces.

Multicast Communication

Broadcast – sends a single message from one process to **all** processes (hosts)

- Used for ARP in a LAN
- Hard and expensive in WAN

Multicast – sends a single message from one process to members of a group of processes (hosts)


Who needs multicast?

Who should provide it?

Application, transport, network layer?

<https://www.youtube.com/watch?v=C5pFaZtbISo>

Multicast Communication



- In multicast routing, the router may forward the received packet through several of its interfaces

Who needs it?

Uses of Multicast and Its Effects

Fault tolerance based on replicated services

- Requests multicast to group of servers

Discovery in spontaneous networking

- Locate available discovery services

Performance from replicated data

- Multicast changes to all replicas

Propagation of event notifications in a distributed environment

- News group: news → group of interested users

Multicast IP address

	oc tet 1	oc tet 2	oc tet 3	Range of addresses
	Network ID		Host ID	
Class A:	1 to 127	0 to 255	0 to 255	1.0.0.0 to 127.255.255.255
	Network ID		Host ID	
Class B:	128 to 191	0 to 255	0 to 255	128.0.0.0 to 191.255.255.255
	Network ID		Host ID	
Class C:	192 to 223	0 to 255	0 to 255	192.0.0.0 to 223.255.255.255
			Host ID	
Class D (multicast):	224 to 239	0 to 255	0 to 255	224.0.0.0 to 239.255.255.255
			Host ID	
Class E (reserved):	240 to 255	0 to 255	0 to 255	240.0.0.0 to 255.255.255.255
			Host ID	

224.0.0.0 to 224.0.0.255 (224.0.0.0/24) → **local** subnet multicast traffic

224.0.1.0 to 238.255.255.255 → **globally** scoped addresses

239.0.0.0 to 239.255.255.255 (239.0.0.0/8) → **administratively** scoped addresses, boundary

IP Multicast Process

Each multicast address → identify a group

Internet Group Membership Protocol (IGMP)

- Processes **register** a group with **local router** using **IGMP**

Router update its multicast routing table

Processes send message to a group

- Do not need to be a member

Router forward multicast messages


Multicast Routing Problem

Goal: find a tree (or trees) connecting routers having local multicast group members


tree: not all paths between routers used

source-based: different tree from each sender to receivers

shared-tree: same tree used by all group members



Shared tree



Source-based trees